

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score, confusion_matrix, mean_absolute_error, mean_squared_error
from sklearn.model_selection import train_test_split
```

```
m = pd.read_csv('med_insurance.csv')
```

```
m.head(10)
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
5	31	female	25.740	0	no	southeast	3756.62160
6	46	female	33.440	1	no	southeast	8240.58960
7	37	female	27.740	3	no	northwest	7281.50560
8	37	male	29.830	2	no	northeast	6406.41070
9	60	female	25.840	0	no	northwest	28923.13692

```
m.isnull()
```

	age	sex	bmi	children	smoker	region	charges
0	False	False	False	False	False	False	False

```
X.head()
```

	age	sex	bmi	children	smoker	region
0	19	female	27.900	0	yes	southwest
1	18	male	33.770	1	no	southeast
2	28	male	33.000	3	no	southeast
3	33	male	22.705	0	no	northwest
4	32	male	28.880	0	no	northwest

```
y.head()
```

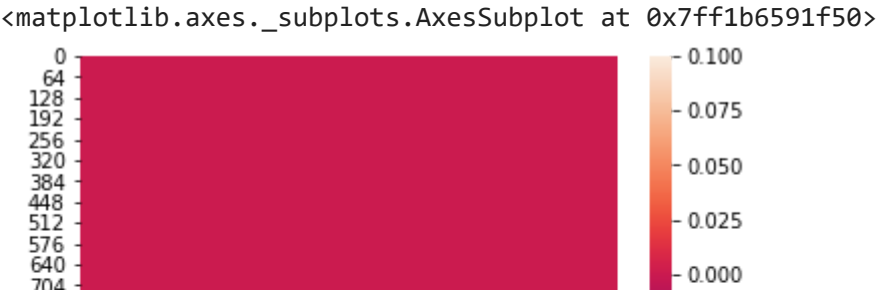
```
0    16884.92400
1     1725.55230
2     4449.46200
3    21984.47061
4     3866.85520
Name: charges, dtype: float64
```

```
m.describe
```

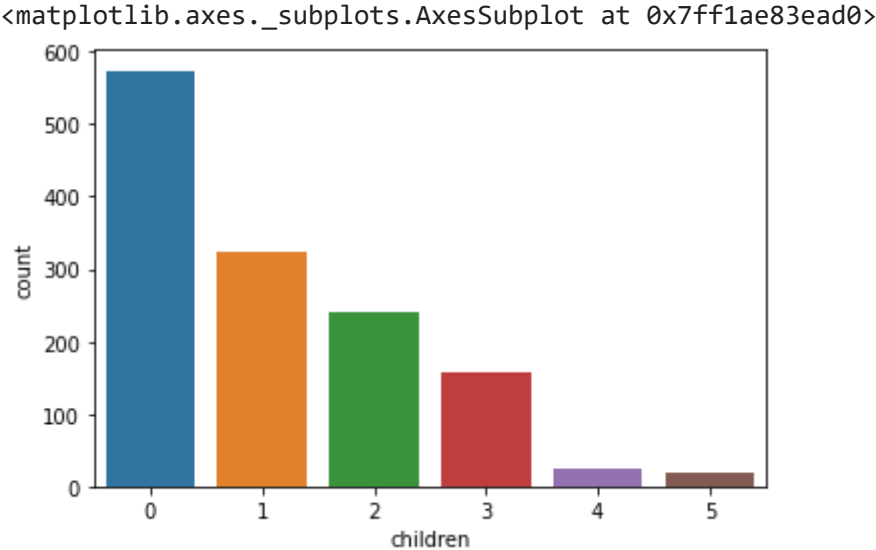
```
<bound method NDFrame.describe of
0    19  female  27.900    0    yes  southwest  16884.92400
1    18   male  33.770    1    no   southeast   1725.55230
2    28   male  33.000    3    no   southeast   4449.46200
3    33   male  22.705    0    no  northwest  21984.47061
4    32   male  28.880    0    no  northwest   3866.85520
...  ...    ...    ...    ...    ...    ...    ...
1333  50   male  30.970    3    no  northwest  10600.54830
1334  18  female  31.920    0    no  northeast   2205.98080
1335  18  female  36.850    0    no  southeast   1629.83350
1336  21  female  25.800    0    no  southwest   2007.94500
1337  61  female  29.070    0   yes  northwest  29141.36030

[1338 rows x 7 columns]>
```

```
sns.heatmap(m.isnull())
```



```
sns.countplot(x='children',data=m)
```



```
g=pd.get_dummies(m['sex'])
```

g

```
female male
r=pd.get_dummies(m['region'])
r
```

	northeast	northwest	southeast	southwest
0	0	0	0	1
1	0	0	1	0
2	0	0	1	0
3	0	1	0	0
4	0	1	0	0
...
1333	0	1	0	0
1334	1	0	0	0
1335	0	0	1	0
1336	0	0	0	1
1337	0	1	0	0

1338 rows × 4 columns

```
s=pd.get_dummies(m['smoker'])
s
```

	no	yes
0	0	1
1	1	0
2	1	0
3	1	0
4	1	0
...
1333	1	0
1334	1	0
1335	1	0
1336	1	0
1337	0	1

1338 rows × 2 columns

```
m=m.drop(['sex','region','smoker'],axis=1)
```

m

	age	bmi	children	charges
0	19	27.900	0	16884.92400
1	18	33.770	1	1725.55230
2	28	33.000	3	4449.46200
3	33	22.705	0	21984.47061
4	32	28.880	0	3866.85520
...
1333	50	30.970	3	10600.54830
1334	18	31.920	0	2205.98080
1335	18	36.850	0	1629.83350
1336	21	25.800	0	2007.94500
1337	61	29.070	0	29141.36030

1338 rows × 4 columns

```
med = pd.concat([g,r,s],axis=1)
med
```

	female	male	northeast	northwest	southeast	southwest	no	yes
0	1	0	0	0	0	1	0	1
1	0	1	0	0	1	0	1	0
2	0	1	0	0	1	0	1	0
3	0	1	0	1	0	0	1	0
4	0	1	0	1	0	0	1	0
...
1333	0	1	0	1	0	0	1	0
1334	1	0	1	0	0	0	1	0
1335	1	0	0	0	1	0	1	0
1336	1	0	0	0	0	1	1	0
1337	1	0	0	1	0	0	0	1

1338 rows × 8 columns

m

	age	bmi	children	charges
0	19	27.900	0	16884.92400
1	18	33.770	1	1725.55230
2	28	33.000	3	4449.46200
3	33	22.705	0	21984.47061
4	32	28.880	0	3866.85520
...
1333	50	30.970	3	10600.54830
1334	18	31.920	0	2205.98080
1335	18	36.850	0	1629.83350
1336	21	25.800	0	2007.94500
1337	61	29.070	0	29141.36030

1338 rows × 4 columns

```
med =pd.concat([m,g,r,s],axis=1)
```

med

	age	bmi	children	charges	female	male	northeast	northwest	southeast
0	19	27.900	0	16884.92400	1	0	0	0	
1	18	33.770	1	1725.55230	0	1	0	0	
2	28	33.000	3	4449.46200	0	1	0	0	
3	33	22.705	0	21984.47061	0	1	0	1	
4	32	28.880	0	3866.85520	0	1	0	1	
...
1333	50	30.970	3	10600.54830	0	1	0	1	
1334	18	31.920	0	2205.98080	1	0	1	0	
1335	18	36.850	0	1629.83350	1	0	0	0	
1336	21	25.800	0	2007.94500	1	0	0	0	
1337	61	29.070	0	29141.36030	1	0	0	1	

1338 rows × 12 columns

```
med.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1338 entries, 0 to 1337
```

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	age	1338 non-null	int64
1	bmi	1338 non-null	float64
2	children	1338 non-null	int64
3	charges	1338 non-null	float64
4	female	1338 non-null	uint8
5	male	1338 non-null	uint8
6	northeast	1338 non-null	uint8
7	northwest	1338 non-null	uint8
8	southeast	1338 non-null	uint8
9	southwest	1338 non-null	uint8
10	no	1338 non-null	uint8
11	yes	1338 non-null	uint8

dtypes: float64(2), int64(2), uint8(8)
memory usage: 52.4 KB

```
y = med.charges
```

```
x = med.drop(['charges'],axis=1)
```

```
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.30,random_state=0)
```

```
regr = LinearRegression()
regr.fit(X_train,y_train)
```

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

```
y_pred = regr.predict(X_test)
```

```
y_pred
```

```
5.55574972e+03, 5.14017616e+03, 6.54358067e+03, 5.29101323e+03,
9.86573296e+03, 5.30318056e+03, 5.82126396e+03, 6.71264320e+03,
3.79359032e+03, 5.35674887e+03, 3.81050149e+04, 1.49449863e+03,
1.25241127e+04, 8.81600344e+03, 1.35779650e+04, 5.33467494e+03,
5.30034294e+03, 3.64048098e+04, 4.39998407e+03, 2.07683907e+03,
1.50700993e+04, 1.28583285e+04, 3.50480278e+04, 4.81692549e+03,
5.75568624e+03, 3.10570003e+04, 5.91365180e+03, 2.12466743e+03,
8.55473411e+03, 1.02157331e+04, 8.01035514e+03, 5.56732969e+03,
1.33173353e+04, 3.83325588e+04, 1.36438797e+04, 2.88073730e+04,
6.57855063e+03, 3.54156777e+04, 3.85264559e+03, 1.18931235e+04,
9.11693680e+03, 6.16245674e+03, 1.11692747e+04, 1.47049487e+04,
5.02634820e+03, 4.46458369e+03, 7.96288703e+03, 1.31978165e+03,
8.00636697e+03, 4.52316428e+03, 1.30510837e+04, 4.39675270e+03,
1.02462007e+04, 7.38951728e+03, 9.36630903e+03, 2.44501178e+03,
1.33392136e+04, 1.66792757e+04, 1.50554402e+04, 1.06897979e+04,
5.36806605e+03, 2.16568099e+03, 1.84625678e+03, 1.36078096e+04,
1.41046244e+04, 5.18518795e+03, 3.78006430e+03, 9.57585167e+03,
1.01219994e+04, 2.81258257e+04, 7.73311859e+03, 1.06662041e+04,
6.36597477e+03, 2.94855951e+04, 1.12667069e+04, 7.60247088e+03,
1.00997356e+04, 1.20691312e+04, 3.08178430e+03, 1.06992857e+04,
1.58012525e+03, 7.13244227e+03, 2.82894803e+04, 3.87009694e+04,
6.44760486e+03, 8.18263663e+03, 2.56156737e+03, 4.43116587e+02,
1.06319373e+04, 4.14413001e+03, 4.80964514e+03, 2.35032467e+03,
```

```

6.90487999e+03, 3.33519011e+04, 3.81889733e+04, 1.46742863e+04,
8.06796404e+03, 1.62426303e+04, 3.28132879e+04, 9.64677964e+03,
3.35687219e+04, 3.28499010e+03, 3.07421002e+04, 8.27403986e+03,
1.43749371e+04, 3.94363238e+03, 3.20677633e+04, 8.50722697e+03,
1.12272940e+04, 9.17836109e+03, 4.02094225e+03, 1.28491704e+04,
1.15880228e+04, 8.28529095e+03, 1.34436097e+04, 2.88016799e+03,
1.07598245e+04, 5.60040795e+03, 1.10443396e+04, 3.15748920e+04,
9.86714924e+03, 1.27147820e+03, 2.86568919e+02, 3.99183572e+04,
9.82112709e+03, 6.98485911e+03, 1.38987602e+04, 1.36056308e+04,
2.72298367e+04, 7.28448164e+03, 6.99440177e+03, 1.22798274e+04,
2.68645584e+03, 3.63656263e+03, 2.49800939e+04, 2.58300744e+04,
1.31031820e+04, 3.32140682e+03, 4.86248281e+03, 9.49806860e+03,
1.26028085e+04, 2.34806136e+04, 3.05185854e+04, 1.02915974e+04,
2.36864639e+04, 2.70214254e+03, 1.13613494e+04, 7.31417605e+03,
8.37333217e+03, -1.20090918e+01, 7.87352474e+03, 3.53845654e+04,
6.34875226e+03, 6.42543828e+03, 2.58973345e+01, 1.06566943e+04,
6.52223275e+03, 9.67955134e+03, 3.90212856e+04, 2.73398516e+04,
1.16690107e+04, 3.52946768e+04, 1.52284127e+04, 6.70599088e+03,

1.07444457e+04, 7.07603187e+03, 3.64305215e+04, 5.77215659e+03,
1.11987869e+04, 9.12757093e+02, 2.39243545e+04, 1.72536528e+03,
3.44370930e+04, 1.11893513e+04, 1.64284856e+03, 3.23084344e+04,
6.83933850e+03, 5.38358717e+03, 3.77008671e+04, 2.38518793e+03,
9.73769497e+03, 2.51677325e+03, 1.29924006e+04, 1.14558043e+03,
1.09157187e+04, 6.82118928e+03, 3.65819756e+04, 7.32055769e+03,
3.03617950e+04, 2.92892602e+04, 6.82053534e+03, 1.09348851e+04,
1.76251405e+03, 2.37139310e+03, 3.67942218e+03, 1.26348376e+04,
3.68242957e+04, 9.88276212e+03, 5.21130967e+02, 1.15304510e+04,
4.96731683e+03, 1.00055528e+04, 5.74380019e+03, 7.07572800e+03,
4.03843476e+03, 2.82570712e+04, 4.44634108e+03, -1.24189730e+03,
3.29513389e+04, 1.26555619e+04, 3.59905000e+04, 1.00050499e+04,
7.57606854e+03, -2.70310110e+02, 2.44881699e+03, 1.17300960e+04,
5.85414692e+03, 3.48550916e+03, 1.22947034e+04, 7.91825176e+03,
7.10941145e+03, 5.62219064e+03, 2.93138277e+03, 3.20691197e+04,
3.42234717e+03, 8.77246198e+03, 4.58681719e+03, 1.32668158e+04,
1.49586949e+04, 7.41278952e+03, 2.66486545e+04, 1.40642726e+04,
1.71400262e+04, 1.16484659e+04])

```

y_test

```

578      9724.53000
610      8547.69130
569     45702.02235
1034    12950.07120
198      9644.25250
...
1261     3277.16100
494     17942.10600
97      10226.28420
418     14418.28040
920     13451.12200

```

Name: charges, Length: 402, dtype: float64

```

weights=regr.coef_
intercept=regr.intercept_
print(weights,intercept)

```

```

[ 256.43544682  335.36907276  472.70978916  23.77337759
 -23.77337759  589.02469054  27.12354518 -405.723989

```


-210.42424672 -11717.99558474 11717.99558474] -730.1121662702444

```
regr.score(X_test,y_test)
```


0.7909160991789905

```
df=pd.DataFrame({'Actual' : y_test,'Predicted' : y_pred})
df
```

	Actual	Predicted
578	9724.53000	11253.193646
610	8547.69130	9544.907094
569	45702.02235	37849.801048
1034	12950.07120	16069.269685
198	9644.25250	6734.408723
...
1261	3277.16100	7412.789524
494	17942.10600	26648.654504
97	10226.28420	14064.272563
418	14418.28040	17140.026157
920	13451.12200	11648.465902

402 rows × 2 columns

```
df1=df.head(25)
df1.plot(kind="bar",figsize=(16,10))
plt.grid(which="major",linestyle='-',linewidth="0.5",color="green")
plt.grid(which="minor",linestyle=':',linewidth="0.5",color="black")
plt.show()
```



```
print('Mean Absolute Error: ',mean_absolute_error(y_test,y_pred))  
print('Mean Square Error: ',mean_squared_error(y_test,y_pred))  
print('Root Mean Absolute Error: ',np.sqrt(mean_squared_error(y_test,y_pred)))
```

```
Mean Absolute Error: 4011.4496793279864  
Mean Square Error: 33342497.826954577  
Root Mean Absolute Error: 5774.296305780867
```

✓ 0s completed at 10:13 AM

